

Diabetes Prediction

Group Members:

Oliver Grudzinski

Pranaav Paladugu

Alec Weinbender

Research Question

With this project, we are attempting to see whether there are any variables that we can use to diagnostically predict whether or not a patient has diabetes, based on the data in the dataset.

Hypotheses

Hypotheses Statements:

Null Hypothesis Example - A given variable does not influence whether a patient may have diabetes.

Alternative Hypothesis Example - A given variable does influence whether a patient may have diabetes.

These are all the variables that we will be looking at: Pregnancies, Glucose, Blood Pressure, Skin Thickness, Insulin, BMI, Diabetes Pedigree Function, Age.

Examples:

H₀ (Glucose) – Glucose levels do not influence whether a patient may have diabetes.

H_A (Glucose) – Glucose levels do have an influence on whether a patient may have diabetes.

H₀ (BMI) – The patient's BMI does not influence whether a patient may have diabetes.

H_A (BMI) – The patient's BMI does influence on whether a patient may have diabetes.

H₀ (Age) – The patient's age does not influence whether a patient may have diabetes.

H_A (Age) – The patient's age does influence on whether a patient may have diabetes.

H₀ (Diabetes Pedigree Function) – The patient's DPF levels do not influence whether a patient may have diabetes.

H_A (Diabetes Pedigree Function) – The patient's DPF levels do have an influence on whether a patient may have diabetes.

Dataset

Explanation of the Dataset:

This dataset was created to do what we are doing, which is to diagnostically predict whether or not a patient has diabetes. All of the data in this dataset is taken from female patients that are at least 21 years old and of Pima Indian heritage.

Source: <https://www.kaggle.com/datasets/uciml/pima-indians-diabetes-database?resource=download>

Description of Features:

Pregnancies: Number of times pregnant

Glucose: Plasma glucose concentration at 2 hours in an oral glucose tolerance test

Blood Pressure: Diastolic blood pressure (mm Hg)

Skin Thickness: Triceps skin fold thickness (mm)

Insulin: 2-Hour serum insulin (mu U/ml)

BMI: Body mass index (weight in kg/(height in m)²)

Diabetes Pedigree Function: Diabetes pedigree function

Age: patient's (years)

Target Variable:

The target variable is Outcomes:

0 = Negative for Diabetes (500)

1 = Positive for Diabetes (268)

```
In [1]: import pandas as pd

df = pd.read_csv("data/diabetes.csv")
df.head()
```

```
Out[1]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Data Preprocessing

Identifying Features

```
In [2]: df.describe()
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.345585
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.477297
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

We are going to move forward with using all of the features given in the dataframe as they are all distinct and have the ability to be used for testing and modeling.

```
In [3]: features = df.columns.drop('Outcome').tolist()
features
```

```
Out[3]: ['Pregnancies',
         'Glucose',
         'BloodPressure',
         'SkinThickness',
         'Insulin',
         'BMI',
         'DiabetesPedigreeFunction',
         'Age']
```

Handling Missing and/or Null Values

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin                768 non-null    int64
5   BMI                   768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                   768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

No null values in the dataframe!

Searching for Miscalculated Zeros

```
In [5]: for column in features:
        num_zeros = (df[column] == 0).sum()
        print(f"Number of 0s in '{column}': {num_zeros}")
```

```
Number of 0s in 'Pregnancies': 111
Number of 0s in 'Glucose': 5
Number of 0s in 'BloodPressure': 35
Number of 0s in 'SkinThickness': 227
Number of 0s in 'Insulin': 374
Number of 0s in 'BMI': 11
Number of 0s in 'DiabetesPedigreeFunction': 0
Number of 0s in 'Age': 0
```

Looking at the columns with 0's, we should replace or drop the values that are continuous variables that don't make sense having zero's.

For this, it makes sense to replace the columns of Glucose, Blood Pressure, Skin Thickness, Insulin, and BMI as they are all variables that should be above zero.

We are going to replace the values by entering the mean values of the columns where there the values are 0.

```
In [6]: columns_to_check = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']

for column in columns_to_check:
    mean_value = df.loc[df[column] != 0, column].mean() # skip 0 values
    df[column] = df[column].replace(0, mean_value)
    num_zeros = (df[column] == 0).sum()
    print(f"Number of 0s in '{column}': {num_zeros}")
```

```
Number of 0s in 'Glucose': 0
Number of 0s in 'BloodPressure': 0
Number of 0s in 'SkinThickness': 0
Number of 0s in 'Insulin': 0
Number of 0s in 'BMI': 0
```

```
In [7]: df.describe()
```

```
Out[7]:
```

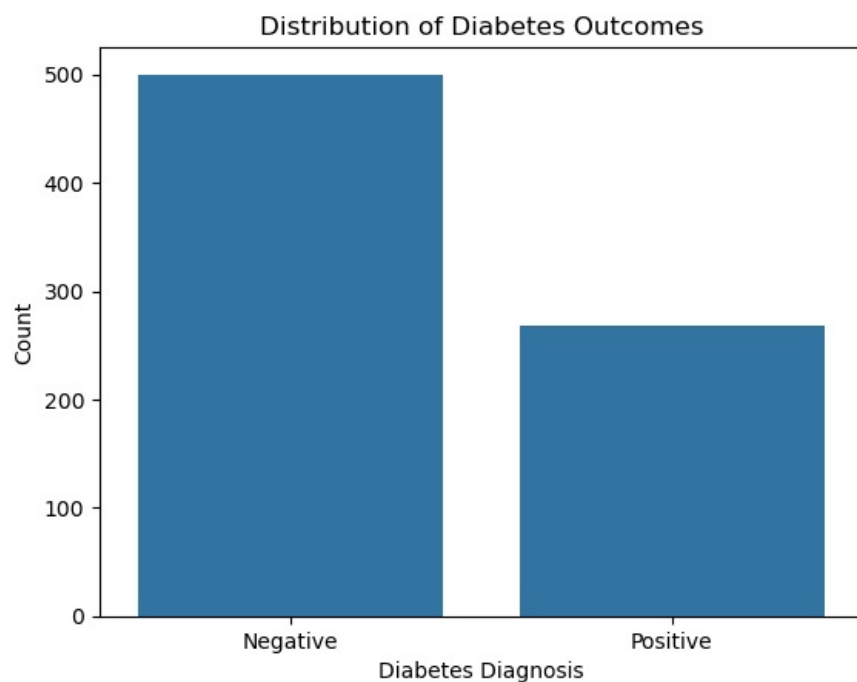
	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	121.686763	72.405184	29.153420	155.548223	32.457464	0.471876	33.240885	0.409090
std	3.369578	30.435949	12.096346	8.790942	85.021108	6.875151	0.331329	11.760232	0.500000
min	0.000000	44.000000	24.000000	7.000000	14.000000	18.200000	0.078000	21.000000	0.000000
25%	1.000000	99.750000	64.000000	25.000000	121.500000	27.500000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.202592	29.153420	155.548223	32.400000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	155.548223	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

Data Analysis and Visualization

```
In [8]: import matplotlib.pyplot as plt
import seaborn as sns
```

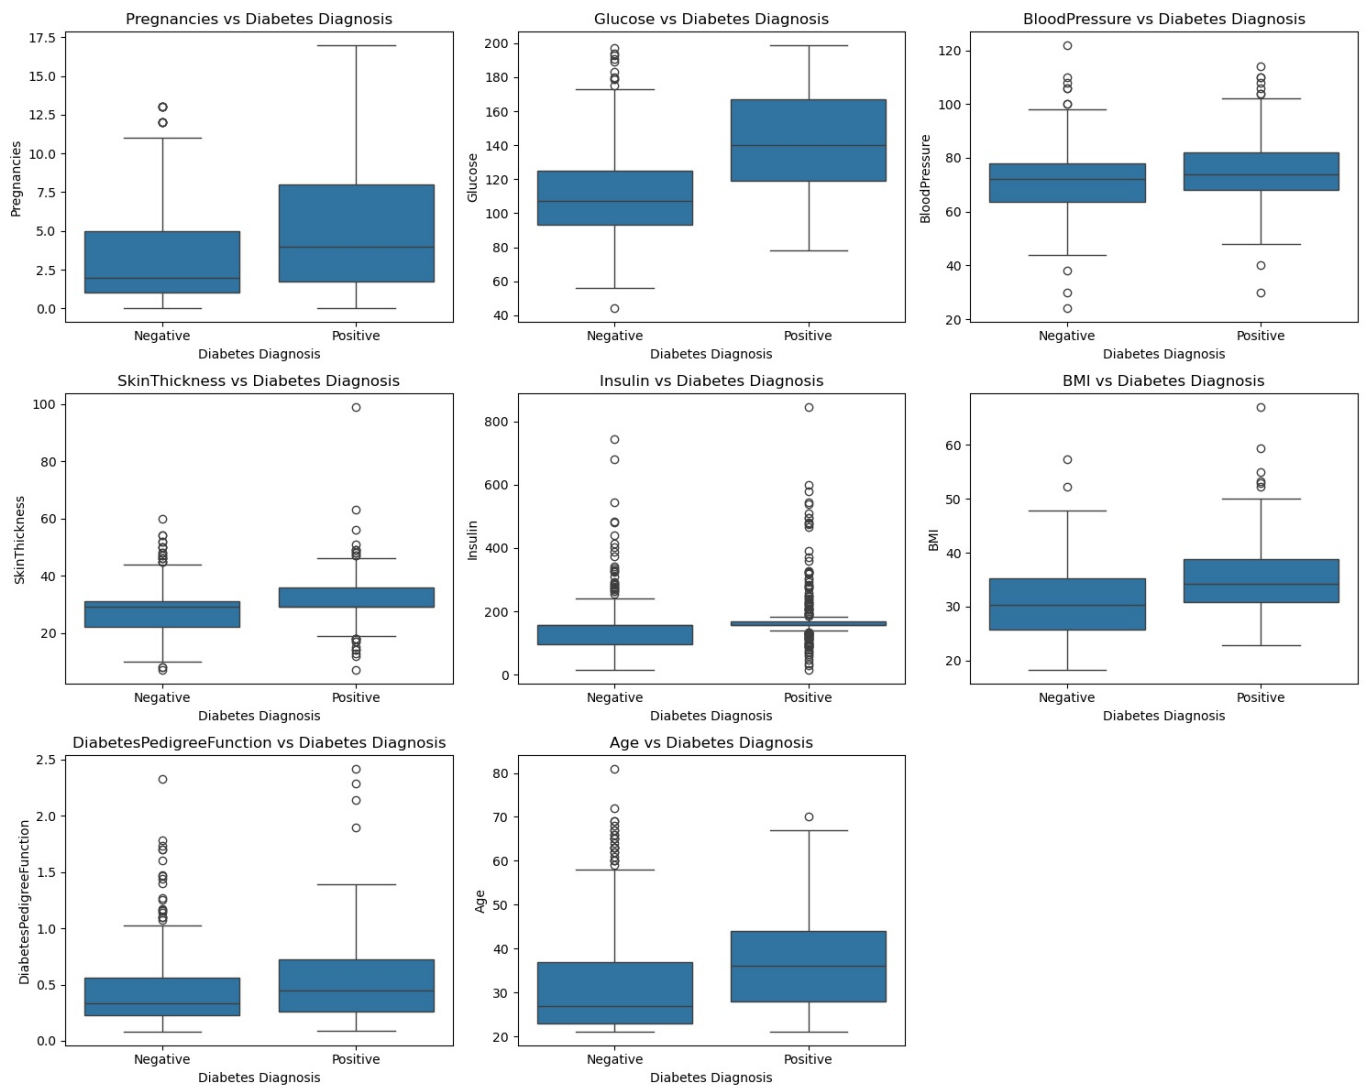
Viewing the Data after Preprocessing

```
In [9]: outcomes = df['Outcome']
sns.countplot(x=outcomes)
plt.xticks([0, 1], ["Negative", "Positive"])
plt.xlabel("Diabetes Diagnosis")
plt.ylabel("Count")
plt.title("Distribution of Diabetes Outcomes")
plt.show()
```



```
In [10]: features = df.columns.drop('Outcome')
num_features = len(features)
cols = 3
rows = (num_features + cols - 1) // cols
plt.figure(figsize=(5 * cols, 4 * rows))
for i, feature in enumerate(features):
    plt.subplot(rows, cols, i + 1)
    sns.boxplot(x=outcomes, y=df[feature])
    plt.xlabel("Diabetes Diagnosis")
    plt.ylabel(feature)
    plt.xticks([0, 1], ["Negative", "Positive"])
    plt.title(f"{feature} vs Diabetes Diagnosis")

plt.tight_layout()
plt.show()
```



Hypothesis Testing

```
In [11]: from scipy import stats
import pandas as pd

target = df["Outcome"]
features = df.columns.drop("Outcome")
n_samples = len(df)

alpha = 0.05
num_tests = len(features)
threshold = alpha/num_tests #Bonferroni Correction

Dict = {}
var = []
pval = []
Dict = {"Variable": [], "T-stat": [], "P-value": [], "Significant": []}
count = 0
for variable in features:
    samples_by_group = []
    for value in set(df[variable]):
        mask = df[variable] == value
        samples_by_group.append(df['Outcome'][mask])
    kstat, Kp = stats.kruskal(*samples_by_group)
    significant = True
    if(Kp < ((threshold)/n_samples)):
        significant = True
    elif(Kp > ((threshold)/n_samples)):
        significant = False
    Dict["T-stat"].append(kstat)
    Dict["P-value"].append(Kp)
    pval.append(Kp)
    var.append(variable)
    Dict["Significant"].append(significant)
    Dict["Variable"].append(variable)

table = pd.DataFrame(Dict)

sorted_df = table.sort_values(by=['P-value'], ascending=True)
```

```
sorted_df.head(10)
```

```
Out[11]:
```

	Variable	T-stat	P-value	Significant
1	Glucose	269.382027	5.586210e-11	True
7	Age	140.754007	2.450833e-10	True
0	Pregnancies	64.510701	8.940856e-08	True
3	SkinThickness	73.467109	1.699207e-02	False
4	Insulin	227.473254	1.818808e-02	False
5	BMI	286.097245	4.422738e-02	False
2	BloodPressure	54.862435	1.738525e-01	False
6	DiabetesPedigreeFunction	532.330054	3.002203e-01	False

Visualizing Significant Variables


```
In [12]: import matplotlib.pyplot as plt
import numpy as np

x = np.array(sorted_df['Variable'])
y = np.array(sorted_df['P-value'])

f = plt.figure()
f.set_figwidth(15)

plt.xlabel('variable')
plt.ylabel('p-value')

plt.scatter(x, y)
plt.hlines(y=threshold, xmin=0, xmax=7, linewidth=2, color='r')
plt.show()
```



```
In [13]: significant_variables = sorted_df[sorted_df["Significant"] == True]["Variable"].tolist()
significant_variables
```

```
Out[13]: ['Glucose', 'Age', 'Pregnancies']
```

We will reject the null hypotheses for the variables Glucose, Age, and Pregnancies.

We have failed to reject the null hypotheses for the variables Skin Thickness, Insulin, BMI, Blood Pressure, and DiabetesPedigreeFunction.

Data Modeling and Prediction

Scaling/Normalizing

We will not need to scale or normalize our data as we are going to be building a Random Forest Classifier model.

Baseline Model with Optimized Parameters

Finding Optimal max_leaf_nodes

```
In [14]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```

import matplotlib.pyplot as plt

# Define features and target variable
X = df[features]
y = df["Outcome"]

X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp)

# Range of max_leaf_nodes to test
leaf_nodes_range = [2, 5, 10, 15, 20, 25, 30, 40, 50, 60, 70, 80, 90, 100, 125, 150]

# Dictionary to store validation accuracies
validation_accuracies = {}

for node_value in leaf_nodes_range:
    model = RandomForestClassifier(random_state=42, n_estimators=50, max_leaf_nodes=node_value)
    model.fit(X_train, y_train)
    y_val_pred = model.predict(X_val)
    accuracy = accuracy_score(y_val, y_val_pred)
    validation_accuracies[node_value] = accuracy
    print(f"max_leaf_nodes: {node_value}, Validation Accuracy: {accuracy:.4f}")

# Find the best max_leaf_nodes from the best validation accuracy
max_leaf_nodes = max(validation_accuracies, key=validation_accuracies.get)
print(f"\nOptimal number of max_leaf_nodes: {max_leaf_nodes}")

# Plot validation accuracies
plt.figure(figsize=(10, 6))
plt.plot(leaf_nodes_range, validation_accuracies.values(), marker='o', linestyle='-.')
plt.title("Validation Accuracy vs. Number of Leaf Nodes")
plt.xlabel("Number of Leaf Nodes")
plt.ylabel("Validation Accuracy")
plt.ylim(0.7, 0.85) # Set y-axis to start at 0.7
plt.grid()
plt.show()

```

```


max_leaf_nodes: 2, Validation Accuracy: 0.6696
max_leaf_nodes: 5, Validation Accuracy: 0.7565
max_leaf_nodes: 10, Validation Accuracy: 0.7652
max_leaf_nodes: 15, Validation Accuracy: 0.7739
max_leaf_nodes: 20, Validation Accuracy: 0.7565
max_leaf_nodes: 25, Validation Accuracy: 0.7826
max_leaf_nodes: 30, Validation Accuracy: 0.7565
max_leaf_nodes: 40, Validation Accuracy: 0.7652
max_leaf_nodes: 50, Validation Accuracy: 0.7652
max_leaf_nodes: 60, Validation Accuracy: 0.7826
max_leaf_nodes: 70, Validation Accuracy: 0.7826
max_leaf_nodes: 80, Validation Accuracy: 0.7826
max_leaf_nodes: 90, Validation Accuracy: 0.7826
max_leaf_nodes: 100, Validation Accuracy: 0.7826
max_leaf_nodes: 125, Validation Accuracy: 0.7826
max_leaf_nodes: 150, Validation Accuracy: 0.7826

```

```

Optimal number of max_leaf_nodes: 25

```



Finding Optimal `n_estimators`

```
In [15]: # Define features and target variable
X = df[features]
y = df["Outcome"]

# Split the data into training, validation, and test sets
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp)

# Range of n_estimators to test
n_estimators_range = [2, 5, 10, 15, 20, 25, 30, 40, 50, 60, 70, 80, 90, 100, 125, 150]

# Dictionary to store validation accuracies
validation_accuracies = {}

# Find the optimal number of estimators
for n_estimators in n_estimators_range:
    model = RandomForestClassifier(random_state=42, n_estimators=n_estimators, max_leaf_nodes=max_leaf_nodes)
    model.fit(X_train, y_train)
    y_val_pred = model.predict(X_val)
    accuracy = accuracy_score(y_val, y_val_pred)
```



```

validation_accuracies[n_estimators] = accuracy
print(f"n_estimators: {n_estimators}, Validation Accuracy: {accuracy:.4f}")

# Find the n_estimators with the best validation accuracy
optimal_n_estimators = max(validation_accuracies, key=validation_accuracies.get)
print(f"\nOptimal number of estimators: {optimal_n_estimators}")

# Plot validation accuracies
plt.figure(figsize=(10, 6))
plt.plot(n_estimators_range, validation_accuracies.values(), marker='o', linestyle='--')
plt.title("Validation Accuracy vs. Number of Estimators")
plt.xlabel("Number of Estimators")
plt.ylabel("Validation Accuracy")
plt.ylim(0.7,0.85) # Set y-axis to start at 0
plt.grid()
plt.show()


```

```

n_estimators: 2, Validation Accuracy: 0.7304
n_estimators: 5, Validation Accuracy: 0.7391
n_estimators: 10, Validation Accuracy: 0.7391
n_estimators: 15, Validation Accuracy: 0.7391
n_estimators: 20, Validation Accuracy: 0.7565
n_estimators: 25, Validation Accuracy: 0.7826
n_estimators: 30, Validation Accuracy: 0.7652
n_estimators: 40, Validation Accuracy: 0.7565
n_estimators: 50, Validation Accuracy: 0.7826
n_estimators: 60, Validation Accuracy: 0.7565
n_estimators: 70, Validation Accuracy: 0.7565
n_estimators: 80, Validation Accuracy: 0.7652
n_estimators: 90, Validation Accuracy: 0.7565
n_estimators: 100, Validation Accuracy: 0.7478
n_estimators: 125, Validation Accuracy: 0.7565
n_estimators: 150, Validation Accuracy: 0.7478

```

Optimal number of estimators: 25



```

In [16]: import matplotlib.pyplot as plt
import pandas as pd
from sklearn.tree import plot_tree

def visualize_forest(forest, feature_names):
    importance_df = pd.DataFrame({
        "Feature": feature_names,
        "Importance": forest.feature_importances_
    }).sort_values("Importance", ascending=False)

    plt.figure(figsize=(10, 6))
    plt.barh(importance_df["Feature"], importance_df["Importance"], color="skyblue")
    plt.gca().invert_yaxis() # Invert y-axis to show the most important feature at the top
    plt.title("Feature Importance in Random Forest Model")
    plt.xlabel("Importance")
    plt.ylabel("Feature")

```

```
plt.tight_layout()
plt.show()

plt.figure(figsize=(20, 10))
plot_tree(forest.estimators_[0], feature_names=feature_names, class_names=["Negative", "Positive"], filled=True)
plt.title("Random Forest Tree Visualization")
plt.show()

X = df[features]
y = df["Outcome"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)


# Train the final model with the optimal number of estimators
optimized_model = RandomForestClassifier(random_state=42, n_estimators=optimal_n_estimators, max_leaf_nodes=max_leaf_nodes)
optimized_model.fit(X_train, y_train)

# Evaluate the model on the test set
optimized_y_test_pred = optimized_model.predict(X_test)
test_accuracy = accuracy_score(y_test, optimized_y_test_pred)
print(f"\nTest Accuracy with {max_leaf_nodes} max_leaf_nodes and {optimal_n_estimators} n_estimators: {test_accuracy}")

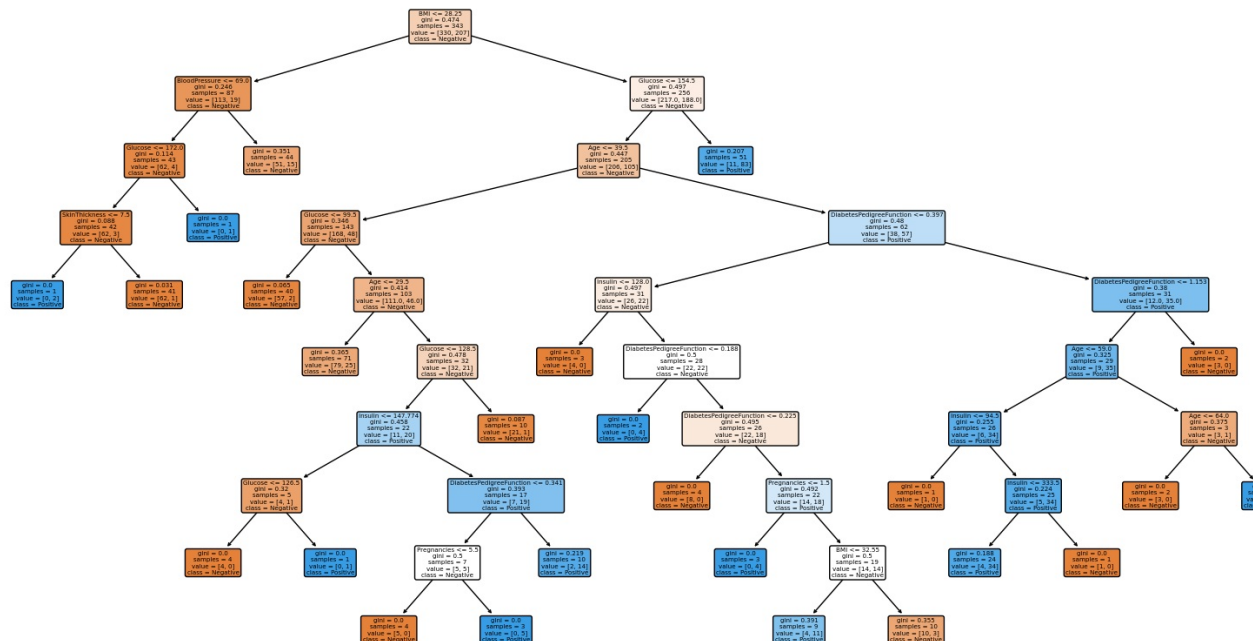
# Feature importance visualization
visualize_forest(optimized_model, features)
```

Test Accuracy with 25 max_leaf_nodes and 25 n_estimators: 0.7706

Feature Importance in Random Forest Model



Random Forest Tree Visualization



Significant Feature Selection Model


```
In [21]: X = df[significant_variables]
y = df["Outcome"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)
significant_model = RandomForestClassifier(random_state=42, n_estimators=optimal_n_estimators, max_leaf_nodes=max_
significant_model.fit(X_train, y_train)

significant_y_test_pred = significant_model.predict(X_test)
test_accuracy = accuracy_score(y_test, significant_y_test_pred)
print(f"\nSignificant Model Test Accuracy with {max_leaf_nodes} max_leaf_nodes and {optimal_n_estimators} n_est

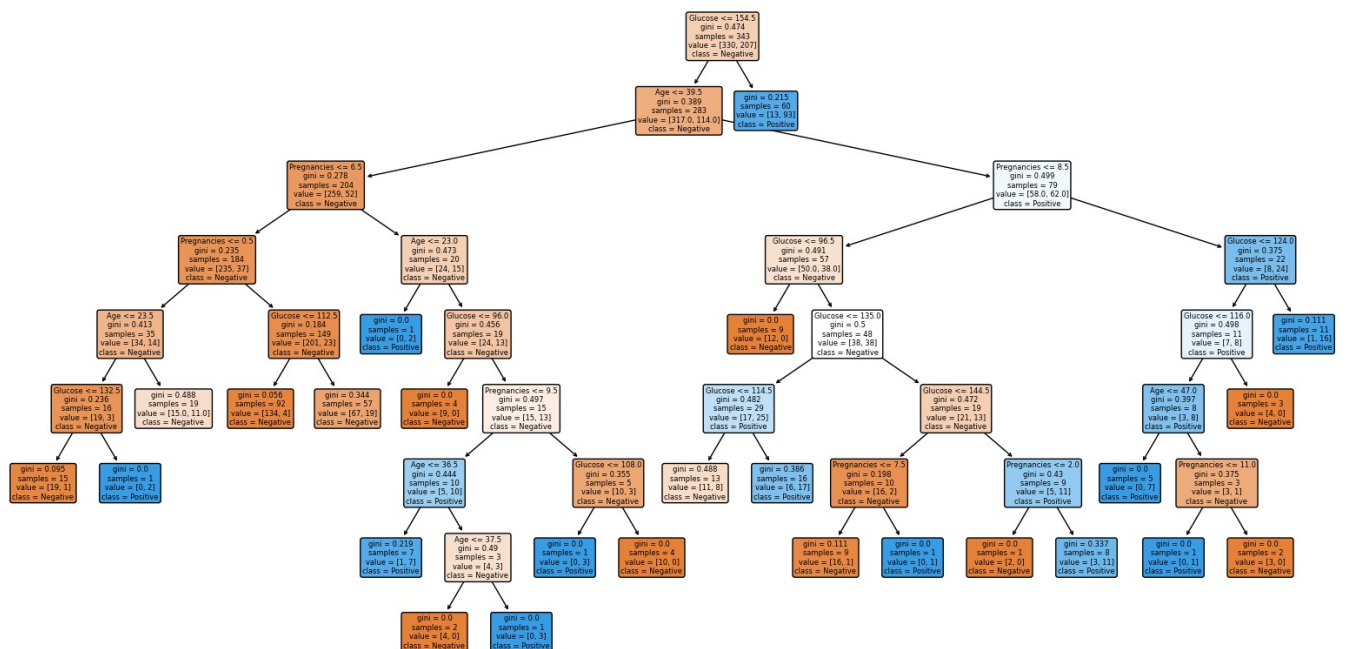
# Feature importance visualization
visualize_forest(significant_model, significant_variables)
```

Significant Model Test Accuracy with 25 max_leaf_nodes and 25 n_estimators: 0.7186

Feature Importance in Random Forest Model



Random Forest Tree Visualization



Greedy Model

```
In [22]: # Rank features by importance
importances = optimized_model.feature_importances_
importances_sorted_indices = importances.argsort()[::-1]
features_sorted = [features[i] for i in importances_sorted_indices]
features_sorted
```

```
Out[22]: ['Glucose',
          'BMI',
          'Age',
          'DiabetesPedigreeFunction',
          'Insulin',
          'Pregnancies',
          'BloodPressure',
          'SkinThickness']
```

```
In [23]: # Add new features to the model one by one (greedy approach)
```

```
greedy_features = []
accuracy = 0
for feature in features_sorted:
    X = df[greedy_features + [feature]]
    y = df["Outcome"]
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)
    model = RandomForestClassifier(random_state=42, n_estimators=optimal_n_estimators, max_leaf_nodes=max_leaf_nodes)
    model.fit(X_train, y_train)

    y_test_pred = model.predict(X_test)
    test_accuracy = accuracy_score(y_test, y_test_pred)
    if test_accuracy > accuracy:
        accuracy = test_accuracy
        greedy_features.append(feature)
        print(f"Added feature: {feature}, Test Accuracy: {test_accuracy:.4f}")
    else:
        print(f"Not added feature: {feature}. Test Accuracy: {test_accuracy:.4f}")

greedy_features
```

```
Added feature: Glucose, Test Accuracy: 0.7013
Added feature: BMI, Test Accuracy: 0.7100
Added feature: Age, Test Accuracy: 0.7446
Added feature: DiabetesPedigreeFunction, Test Accuracy: 0.7706
Not added feature: Insulin. Test Accuracy: 0.7532
Added feature: Pregnancies, Test Accuracy: 0.7749
Not added feature: BloodPressure. Test Accuracy: 0.7619
Not added feature: SkinThickness. Test Accuracy: 0.7532
```

```
Out[23]: ['Glucose', 'BMI', 'Age', 'DiabetesPedigreeFunction', 'Pregnancies']
```


```
In [24]: X = df[greedy_features]
y = df["Outcome"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)

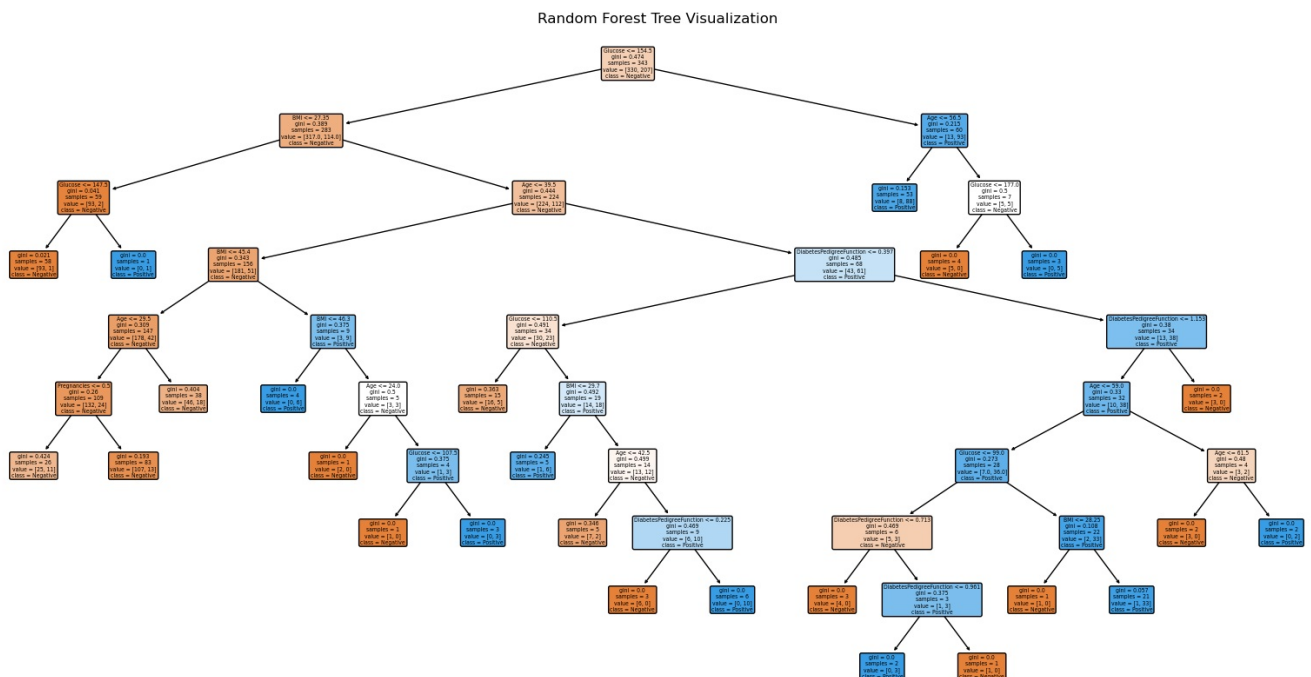
greedy_model = RandomForestClassifier(random_state=42, n_estimators=optimal_n_estimators, max_leaf_nodes=max_leaf_nodes)
greedy_model.fit(X_train, y_train)

greedy_y_test_pred = greedy_model.predict(X_test)
test_accuracy = accuracy_score(y_test, greedy_y_test_pred)
print(f"\nGreedy Model Test Accuracy with {max_leaf_nodes} max_leaf_nodes and {optimal_n_estimators} n_estimators: {test_accuracy:.4f}")

# Feature importance visualization
visualize_forest(greedy_model, greedy_features)
```

```
Greedy Model Test Accuracy with 25 max_leaf_nodes and 25 n_estimators: 0.7532
```





Model Accuracies

```
In [25]: # F1 Score
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score, recall_score
test_predictions = {
    "Baseline Model": optimized_y_test_pred,
    "Significant Model": significant_y_test_pred,
    "Greedy Model": greedy_y_test_pred
}

f1_scores = {}
for model_name, predictions in test_predictions.items():
    f1 = f1_score(y_test, predictions)
    f1_scores[model_name] = f1
    print(f"{model_name} F1 Score: {f1:.4f}")

# Precision and Recall

precision_scores = {}
recall_scores = {}

for model_name, predictions in test_predictions.items():
    precision = precision_score(y_test, predictions)
    recall = recall_score(y_test, predictions)
    precision_scores[model_name] = precision
    recall_scores[model_name] = recall
    print(f"{model_name} Precision: {precision:.4f}")
    print(f"{model_name} Recall: {recall:.4f}")
```

Baseline Model F1 Score: 0.6241
Baseline Model Precision: 0.7333
Baseline Model Recall: 0.5432
Significant Model Precision: 0.6212
Significant Model Recall: 0.5062
Greedy Model Precision: 0.7101
Greedy Model Recall: 0.6049
Significant Model F1 Score: 0.5578
Baseline Model Precision: 0.7333
Baseline Model Recall: 0.5432
Significant Model Precision: 0.6212
Significant Model Recall: 0.5062
Greedy Model Precision: 0.7101
Greedy Model Recall: 0.6049
Greedy Model F1 Score: 0.6533
Baseline Model Precision: 0.7333
Baseline Model Recall: 0.5432
Significant Model Precision: 0.6212
Significant Model Recall: 0.5062
Greedy Model Precision: 0.7101
Greedy Model Recall: 0.6049

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js